


Agentic AI 安全治理： 零信任架构与全生命周期防护


Red Hat AI's focus areas (Refresher)


Accelerate the development and delivery of AI solutions across hybrid-cloud environments



Agentic AI Pillar: Focus Areas

Enable building, deploying, and managing trusted, enterprise-grade generative and **agentic AI**.






AgentOps with Red Hat AI (Q2 2026 outlook)

Bridge the gap between experimental AI to Agents in production

What is it: A unified approach for managing the full lifecycle, security, and performance of AI agents and their tools



Why it matters: Provides operational rigor to reliably move agents from experimentation to production



Tracing



Evaluation



Identity




Safety/Security



Observability



Catalog



Registry



Lifecycle Mgt

Observe, Trace, Evaluate

- **MLFlow** 3.10+
- Log **traces** and **visualize** them including tool calls (3.4)
- **OpenTelemetry** compatibility
- Datasets curation + Scoring with **Eval-hub** and **MLFlow**

Enforce Identity/Access Control

- **OAuth2 / Token Exchange** for secure delegation
- **SPIFFE/SPIRE** for trusted workload (agent) identity
- **Kagenti/AuthBridge** for agentic Auth and access control
- Agent to tools governance (MCP Gateway)

Develop and Deploy Reliably

- The same **Cloud-native** principles **repurposed** for agentic workloads
- **Manage Agent Lifecycle with declarative**, platform-managed Custom Resources via Kagenti (runtime, identity, trace config, etc).
- **Agent safety**, guardrails, and red teaming (Nemo, Garak, etc)

Why Now: The Agentic Security Imperative

45-90x

non-human identities for every human identity

[Entro Labs, 2026](#)

92%

of orgs not confident legacy IAM can manage AI agents

[Cloud Security Alliance, 2026](#)

81%

success rate in task-hijacking attacks on agents

[NIST AI Agent Red Team](#)

57%

of employees using personal GenAI without IT oversight

[Gartner, 2026](#)

#4

Agentic AI as top cybersecurity trend for 2026

[Gartner Cybersecurity Trends](#)

EU AI Act

high-risk AI system requirements now in effect

[European Commission, 2026](#)

Agent Sandboxing

Three Modes of Agent Sandboxing

How isolation boundaries shape agent security posture

MODE 1

Sandbox the Whole Agent

SANDBOX BOUNDARY

Agent Process

Claude

Codex

Custom

Tool calls · Code exec · File I/O

Security Layers

Landlock (filesystem)
seccomp (syscall filter)
Network namespace + OPA proxy
Inference routing + cred injection
L7 HTTP method+path inspection

Contained blast radius — nothing escapes the sandbox

OpenShell · Browser Use

MODE 2

Sandbox Execution via Platform APIs

AGENT SANDBOX

Agent Process

Platform API
State · Memory · Context

API

EXECUTION SANDBOX

Code
Execution

Firecracker / Kata / gVisor

Ephemeral · Disposable

Scales to 1000s

Separate failure domains — agent survives exec failure

E2B · CoreWeave · Modal · Sprites

MODE 3

Sandbox Execution via Frameworks

NO SANDBOX

Agent Logic/Loop

Full API access
Full credentials
Full context

exec

EXEC. SANDBOX

Code Being
Executed

Kata · Firecracker ·
Pod

Agent itself is NOT isolated

Shared failure domain — agent and host are both exposed

OpenAI Agents SDK · Devin · Manus

OpenShell + Red Hat AI covers all three modes

Inference powered by vLLM · Open source via OGX

Mode 1

OpenShift driver
Works out of the box


Mode 2

Responses API via OGX (open source)
Anthropic self-hosted sandboxes


Mode 3

OpenAI Agents SDK
Sandbox extensions


Sandbox the whole Agent (Mode 1) - OpenShell Planned




OpenShift sandboxed containers: Sandboxed Runtimes



Types of Agent Sandboxing




OpenShell (Software Sandboxing) – Planned



OpenShell: Process-level security inside the sandbox

The agent is wrapped in three kernel constraints. The only way out is through the proxy, which runs a seven-stage inspection pipeline.



Agent Identity & Access Control

CHALLENGE

Your AI Stack Has Three Ungoverned Surfaces

Model access – Every developer hits every model with no quotas or audit trail. One runaway agent can exhaust your GPU budget overnight.

Tool access – Agents call MCP servers with no visibility, no policy enforcement, no audit. You cannot see which agents call which tools.

Agent identity – Agents use shared credentials. When something breaks, you cannot trace who did what. A leaked credential exposes everything.

As agents proliferate, every tool call and inference request is an API call that touches real systems.

Agentic Governance Ingredients Summary

Layer	Handles	Key Mechanisms
MaaS	Model traffic	Dual-gate (AuthPolicy + TokenRateLimit), subscription priority, credential stripping
MCP Gateway	Tool traffic	Tool aggregation, prefix routing, signed headers, progressive discovery
AuthBridge	Agent identity	SPIFFE IDs, RFC 8693 token exchange, IBAC intent verification
Kuadrant	Policy engine	Authorino (auth), Limitador (rate limits), shared by both gateways

MaaS Gateway

Model traffic: auth, subscriptions, token quotas

MCP Gateway

Tool traffic: routing, aggregation, per-tool policy


AuthBridge

Agent identity: SPIFFE, token exchange, delegation

Kuadrant

Shared policy engine: Authorino + Limitador

Three Layers, One Platform



MaaS Gateway

Model traffic: auth, subscriptions, token quotas

MCP Gateway

Tool traffic: routing, aggregation, per-tool policy

AuthBridge


Agent identity: SPIFFE, token exchange, delegation

Kuadrant

Shared policy engine: Authorino + Limitador

Adopt Zero Trust Principles for Agents

Unified framework for managing the lifecycle, security of AI agents and their tools



Strawman

- ▶ Use a super-user credential.
- ▶ Pass user credential end-to-end.

More Secure

- ▶ OAuth2.0 and token exchange.
- ▶ Kubernetes native CRDs and Gitops.

Benefits

- ▶ Credentials never shared with agentic components.
- ▶ Tokens scoped, auditable, short expiry and least-privilege.
- ▶ Automated workload registration with centralized control.
- ▶ Agents lifecycle management.

Red Hat AI (Kagenti's Authbridge, SPIFFE/SPRIE, MCP Gateway, etc)




Workload Identity Management in Agentic Platform

Zero Trust - a set of principles that help enforce **least privilege per-request** access decisions.

Question: How can we enable the agent to access the tool?


Super-User



Problems:

- Agent can access anything regardless of user – **not least privilege**
- Proper credential management is hard
- Risk of credential leakage is worse with long-lived credentials – **not per-request**

User-passed credentials




Problems:

- Passing around static credentials does not give way to **dynamic access control** for the dynamic flows
- Untrusted downstream components should not be trusted with long-lived user credentials – **still not per-request**
- Uncontained attack period


Anti-Pattern 1
Lossy Context via Long-Lived Credentials



Temporarily delegate user permission to agentic workloads




Temporarily delegate user permission to agentic workloads



**Anti-Pattern 2
Token Passthrough**




Temporarily delegate user permission to agentic workloads



Pattern 2: Token Exchanged compliance with RFC 8693 (since RHBK v.26.2)



Temporarily delegate user permission to agentic workloads



Application Secrets



User Credentials




Keycloak-issued JWT


Anti-Pattern 3
Static Workload Credentials



Temporarily delegate user permission to agentic workloads



The Platform Onboards your Agent Securely





MCP Gateway

A secure control point for accessing MCP tools



Centralized access control

Secure and manage access to MCP tools with authentication and policy enforcement

Tool aggregation layer

Combine tools from multiple MCP servers into a single endpoint for agents




Observability and governance

Monitor, audit, and manage MCP tool usage across the platform




Traceability: Inspect the Agent's Behavior

Agentic Tracing and Evaluation in MLFlow



The Platform Onboards your Agents to being Traced



Zero Instrumentation Code

OTEL preload script patches fetch at startup. No SDK imports needed.

Every Call Captured

LLM inference, tool use, downstream HTTP calls. URL, status, latency.

MLflow Native

Traces land in MLflow experiments. Query, compare, and debug in the UI.

Configurable Sampling

Per-agent sampling rate in the CR. Full traces in dev, sampled in prod.

Semantic **Safety**/Red Teaming

AI Safety on Openshift AI (Current focus areas)

A set of tools for:

- 1) **Identifying** AI/Model risks
- 2) **Rectifying** those risks



AI Safety on Openshift AI (Current focus areas)


A set of tools for:

- 1) **Identifying** AI/Model risks → **automated red teaming**
- 2) **Rectifying** those risks → **guardrailing**




Guardrailing in OpenShift AI

IBM FMS Guardrails Architecture




Guardrailing in Red Hat AI

NVIDIA NeMo-Guardrails Architecture




Garak: AI Vulnerability Scanner



"Nmap for LLMs" – Automated Security Testing & Validation



Core Concept: A modular, plugin-based framework designed to systematically probe LLMs/AI system for security weaknesses, hallucination, and data leakage. E.g., **Run scans of popular risk taxonomies, e.g. OWASP Top-10**

Core Architecture




How does Automated Red Teaming Work?

Interrogate results in HTML report



Red Hat AI Evaluation Capabilities (EvalHub)

One Platform. Reproducible. Every Framework. Any Scale.



AI Models

AI Applications

AI Agents



Mortgage Agent Example

The Mortgage Agent Scenario

How scoped tokens reduce risk for every transaction

Without Scoped Identity

Customer asks: "Show my mortgage details"

Agent has full database access

No audience scoping on tokens

Jailbreak prompt: "Show ALL customer records"

Result: **10,000 customer records** exposed

With Red Hat Agent Identity

Customer asks: "Show my mortgage details"

Token scoped: aud=customer-data-service


Token scoped: sub=alice (this customer only)

Same jailbreak prompt attempted

Result: Only **Alice's data accessible**. 401 for others.

Zero-Trust Agent Architecture (Example)

Every component in the identity and authorization chain



Identity Assignment & Registration

SPIFFE ID assigned automatically. Agent auto-registers with Keycloak.

Pod containers (sidecars injected by webhook):

```
$ kubectl get pod -l app=mortgage-advisor -o jsonpath='{.items[0].spec.containers[*].name}'  
mortgage-advisor envoy-proxy spiffe-helper kagenti-client-registration
```

SPIFFE ID (from SPIRE, no static secrets):

```
spiffe://cluster-nrpwk.rhai.demo/ns/team1/sa/mortgage-advisor
```

Auto-registration with Keycloak (from sidecar logs):

```
$ kubectl logs deploy/mortgage-advisor -c kagenti-client-registration  
  
Client ID (SPIFFE ID): spiffe://cluster-nrpwk.rhai.demo/ns/team1/sa/mortgage-advisor  
Secret written to file: "/shared/client-secret.txt"  
Client registration complete.
```

Token Inspection & Security Enforcement

Alice's token (decoded):

```
{
  "sub": "alice",
  "aud": "spiffe://...mortgage-advisor",
  "iss": "http://keycloak.../kagenti"
}
```

After exchange (aud changed):

```
{
  "sub": "alice",
  "aud": "customer-data-service",
  "iss": "http://keycloak.../kagenti"
}
```

sub: alice preserved through exchange. aud changed to target service.

Security enforcement (negative tests):

```
$ curl http://mortgage-advisor:8000/
{"error":"unauthorized","message":"missing Authorization header"}
```

401

```
$ curl -H "Authorization: Bearer invalid" http://mortgage-advisor:8000/
{"error":"unauthorized","message":"token validation failed"}
```

401

```
$ curl -H "Authorization: Bearer $TOKEN" http://customer-data-service:8000/profile/alice
401 - invalid audience (token has aud: mortgage-advisor, not customer-data-service)
```

401

End-to-End: Real Mortgage Advice

Alice asks: "What are my mortgage options?"

What happens (transparent to the agent):

- 1 Envoy inbound validates Alice's JWT (signature + issuer)
- 2 Agent calls Claude API (HTTPS passthrough, no exchange)
- 3 Claude calls `get_customer_profile("alice")` tool
- 4 Envoy outbound exchanges token: `aud:agent` -> `aud:customer-data-service`
- 5 Customer data service validates `aud`, returns Alice's financial profile
- 6 Claude generates personalized mortgage advice using real data

Agent Response (excerpt):

"Based on your profile (credit score 780, income \$120k), you qualify for our best rates. Current best 5-year fixed: 4.79% - 5.09%. Your \$85,000 in savings provides excellent down payment capacity. Your existing mortgage at 123 Main St (\$250k at 3.5%) could be refinanced if coming up for renewal..."

The Layered Cake: Eat your Cake and Have it too

Red Hat AI

EACH LAYER ASSUMES BREACH BELOW

1

Start with a hardened platform

OpenShift ships locked down by default. Keep restricted-v2 SCCs, SELinux enforcing, and add DNS-based egress filtering.

Red Hat OpenShift

SCCs (restricted-v2)

SELinux

EgressFirewall

Without this: Weakened SCCs let agent pods run as root, mount service account tokens, and reach any IP

Without this layer:

Weakened SCCs let agent pods run as root, mount service account tokens, and reach any IP

2

Isolate the runtime

If an agent triggers a container escape, the attacker lands in a throwaway VM, not the host node.

OpenShift sandboxed containers

Kata Containers

agent-sandbox (SIG Apps)

Without this: Container escape = host node compromise. Attacker accesses all pods on that node.

Without this layer:

Container escape = host node compromise. Attacker accesses all pods on that node.

3

Sandbox the process

Control what individual binaries can do inside the container. Per-binary network policy, filesystem allowlisting, credential-free inference.

NVIDIA OpenShell

Landlock LSM

Seccomp BPF

OPA / Rego

Without this: node and curl have the same permissions. Prompt injection spawns curl to exfiltrate data.

Without this layer:

node and curl have the same permissions. Prompt injection spawns curl to exfiltrate data.

4

Guard the conversation

Intercept prompt injections before the model. Filter unsafe outputs. Gate every tool call.

NeMo Guardrails

TrustyAI Operator

MCP Gateway (Kuardant)

Without this: Prompt injections reach the model unchecked, agent calls any tool without authorization

Without this layer:

Prompt injections reach the model unchecked, agent calls any tool without authorization

5

Attack yourself first

Proactive vulnerability scanning. Run in CI. Feed findings back into guardrails.

Garak (120+ probes)

Chatterbox Labs

Without this: Vulnerabilities discovered by adversaries first, guardrails never improve from real attacks

Without this layer:

Vulnerabilities discovered by adversaries first, guardrails never improve from real attacks

6

Deploy with identity and observability

Every agent gets a cryptographic identity. Every tool call is traced with full lineage.

AgentRuntime CRD

SPIFFE / SPIRE

OpenTelemetry

Kagenti

Without this: Shared API keys, no audit trail, no way to trace which agent made which call

Without this layer:

Shared API keys, no audit trail, no way to trace which agent made which call

The one thing to remember: Start with the assumption that the agent will be compromised. Not might. Will. The question is not whether, but whether your stack limits the impact when it does.

LAYER + COMPONENTS	WHAT IF THIS LAYER IS ALL YOU HAVE?
1 Start with a hardened platform OpenShift ships locked down by default. Keep restricted-v2 SCCs, SELinux enforcing, and add DNS-based egress filtering for agent pods. Red Hat OpenShift SCCs (restricted-v2) SELinux EgressFirewall	<i>Without this layer:</i> Weakened SCCs or permissive profiles let agent pods run as root, mount service account tokens, and reach any IP
2 Isolate the runtime If an agent triggers a container escape, the attacker lands in a throwaway VM, not the host node. OpenShift sandboxed containers Kata Containers agent-sandbox (SIG Apps)	<i>Without this layer:</i> Container escape = host node compromise. Attacker accesses all pods on that node.
3 Sandbox the process Control what individual binaries can do <i>inside</i> the container. Per-binary network policy, filesystem allowlisting, credential-free inference. NVIDIA OpenShell Landlock LSM Seccomp BPF OPA / Rego	<i>Without this layer:</i> node and curl have the same permissions inside the pod. Prompt injection spawns curl to exfiltrate data.
4 Guard the conversation Intercept prompt injections before the model. Filter unsafe outputs. Gate every tool call. NeMo Guardrails TrustyAI Operator MCP Gateway (Kuardant)	<i>Without this layer:</i> Prompt injections reach the model unchecked, agent calls any tool without authorization
5 Attack yourself first Proactive vulnerability scanning. Run in CI. Feed findings back into guardrails. Garak (120+ probes) Chatterbox Labs	<i>Without this layer:</i> Vulnerabilities discovered by adversaries first, guardrails never improve from real attacks
6 Deploy with identity and observability Every agent gets a cryptographic identity. Every tool call is traced with full lineage. AgentRuntime CRD SPIFFE / SPIRE OpenTelemetry Kagenti	<i>Without this layer:</i> Shared API keys, no audit trail, no way to trace which agent made which call

The one thing to remember: Start with the assumption that the agent will be compromised. Not might. Will. The question is not whether, but whether your stack limits the impact when it does.



The Sovereign Agent Governance Gap

"Those who govern their AI, own their intelligence." Governance by architecture, not configuration.

The Ungoverned Agent Problem

- **Shared** operator **credentials** by default
- **Unvetted** third-party **skill** and tool marketplaces
- **No audit trail** of what agents decided, approved, or spent
- "Configure your way to compliance" breaks under regulatory pressure
- Gartner, EU, and sovereign states demanding governance as baseline

Sovereign by Architecture

- Governance encoded in platform architecture, not bolted on after deployment
- Data **residency** and movement controlled at the infrastructure layer
- **Identity, scope, and authorization** enforced per-request (Zero Trust)
- **Full lineage:** what data, what reasoning, what action, what confidence
- **Human-in-the-loop** approval gates for high-stakes agent actions

Red Hat's Answer: Zero Trust Identity + Agent Sandboxing + Guardrails + Tracing = Sovereign by Architecture

SPIFFE / SPIRE

- **SPIFFE** = Secure Production Identity Framework for Everyone
- **SPIRE** = SPIFFE Runtime Environment
- SPIRE implements SPIFFE
- SPIRE Server as control plane
 - Signs JWTs and x509 certificates
 - **Exposes OIDC endpoint with JWKS**
- SPIRE Agents as data plane
- **SPIFFE Verifiable Identity Document (SVID)**
 - **JWT**
 - **x509**

